

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified for ECE C147/
C247 at UCLA.
"""

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)
        - y is a numpy array of size (num_examples, )
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2

        num_test = X.shape[0]
        num_train = self.X_train.shape[0]
        dists = np.zeros((num_test, num_train))
        for i in np.arange(num_test):
            for j in np.arange(num_train):
                # ===== #
                # YOUR CODE HERE:
                #   Compute the distance between the ith test point and the jth
                #   training point using norm(), and store the result in dists[i, j].
                # ===== #
                dists[i, j] = norm(self.X_train[j] - X[i])

                # ===== #
                # END YOUR CODE HERE
                # ===== #

        return dists

    def compute_L2_distances_vectorized(self, X):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train WITHOUT using any for loops.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.

```

```

"""
num_test = X.shape[0]
num_train = self.X_train.shape[0]
dists = np.zeros((num_test, num_train))

# ===== #
# YOUR CODE HERE:
# Compute the L2 distance between the ith test point and the jth
# training point and store the result in dists[i, j]. You may
# NOT use a for loop (or list comprehension). You may only use
# numpy operations.
#
# HINT: use broadcasting. If you have a shape (N,1) array and
# a shape (M,) array, adding them together produces a shape (N, M)
# array.
# ===== #

# Let A = X, B = self.X_train
# (A-B)**2 = A**2 - 2AB + B**2

# Have: A = (500, 3072), B = (5000, 3072)
# Want: A**2 = (500, 1), B**2 = (5000,), AB = (500, 5000)

# (500,) array with each row being the sum of a testing data point's coordinates
a_squared = np.sum(np.square(X), axis=1)
# Reshape to (500,1) array
a_squared = a_squared.reshape((a_squared.shape[0], 1))
# (5000,) array with each row being the sum of a test data point's coordinates squared
b_squared = np.sum(np.square(self.X_train), axis=1)
# (500,1) array + (5000,) array = (500, 5000) array that holds A**2 + B**2 for each
# pair of testing points and training points
dists = a_squared + b_squared

# (500, 3072).T * (5000, 3072) = (500, 3072) * (3072, 5000) = (500, 5000)
a_times_b = np.dot(X, self.X_train.T)
dists -= 2*a_times_b
dists = np.sqrt(dists)

# ===== #
# END YOUR CODE HERE
# ===== #

return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      gives the distance between the ith test point and the jth training point.

    Returns:
    - y: A numpy array of shape (num_test,) containing predicted labels for the
      test data, where y[i] is the predicted label for the test point X[i].
    """
    num_test = dists.shape[0]
    y_pred = np.zeros(num_test)
    for i in np.arange(num_test):
        # A list of length k storing the labels of the k nearest neighbors to
        # the ith test point.
        closest_y = []
        # ===== #
        # YOUR CODE HERE:
        # Use the distances to calculate and then store the labels of
        # the k-nearest neighbors to the ith test point. The function
        # numpy.argsort may be useful.
        #
        # After doing this, find the most common label of the k-nearest
        # neighbors. Store the predicted label of the ith training example

```

```

        # as y_pred[i]. Break ties by choosing the smaller label.
        # ===== #

        closest_idx = np.argsort(dists[i])[:k]
        closest_y = [self.y_train[idx] for idx in closest_idx]
        unique, counts = np.unique(closest_y, return_counts=True)
        closest_y = sorted(dict(zip(unique, counts)).items(), key=lambda x: (-x[1], x[0]))
        y_pred[i] = closest_y[0][0]

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    return y_pred

```